

Appendix C

Program Codes

C.1. Parallel BFS

```
/*
Implementing Breadth first search on CUDA based on the module provided by
Rodinia benchmark.
*/
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <math.h>
#include <cuda.h>

#define MAX_THREADS_PER_BLOCK 512

int no_of_nodes;
int edge_list_size;
FILE *fp;

//Structure to hold a node information
struct Node
{
    int starting;
    int no_of_edges;
};

__global__ void
Kernel( Node* g_graph_nodes, int* g_graph_edges, bool* g_graph_mask, bool*
g_updating_graph_mask, bool *g_graph_visited, int* g_cost, int no_of_nodes)
{
    int tid = blockIdx.x*MAX_THREADS_PER_BLOCK + threadIdx.x;
    if( tid<no_of_nodes && g_graph_mask[tid])
    {
        g_graph_mask[tid]=false;
        for(int i=g_graph_nodes[tid].starting;
i<(g_graph_nodes[tid].no_of_edges + g_graph_nodes[tid].starting); i++)
        {
            int id = g_graph_edges[i];
            if(!g_graph_visited[id])
            {
                g_cost[id]=g_cost[tid]+1;
                g_updating_graph_mask[id]=true;
            }
        }
    }
}
```

```

__global__ void
Kernel2( bool* g_graph_mask, bool *g_updating_graph_mask, bool* g_graph_visited,
bool *g_over, int no_of_nodes)
{
    int tid = blockIdx.x*MAX_THREADS_PER_BLOCK + threadIdx.x;
    if( tid<no_of_nodes && g_updating_graph_mask[tid])
    {
        g_graph_mask[tid]=true;
        g_graph_visited[tid]=true;
        *g_over=true;
        g_updating_graph_mask[tid]=false;
    }
}

////////////////////////////////////
//Apply BFS using CUDA
////////////////////////////////////
void BFS()
{
    char *input_f;
    //input_f = "B_graph4096.txt"; //argv[1];
    input_f = (char*)malloc(sizeof(char)*100);
    printf("File path:");
    scanf("%s",input_f);
    printf("Reading File\n");
    //Read in Graph from a file
    fp = fopen(input_f,"r");
    if(!fp)
    {
        printf("Error Reading graph file\n");
        return;
    }

    int source = 0;

    fscanf(fp,"%d",&no_of_nodes);

    int num_of_blocks = 1;
    int num_of_threads_per_block = no_of_nodes;

    //Make execution Parameters according to the number of nodes
    //Distribute threads across multiple Blocks if necessary
    if(no_of_nodes>MAX_THREADS_PER_BLOCK)
    {
        num_of_blocks =
(int)ceil(no_of_nodes/(double)MAX_THREADS_PER_BLOCK);
        num_of_threads_per_block = MAX_THREADS_PER_BLOCK;
    }
}

```

```

}

// allocate host memory
Node* h_graph_nodes = (Node*) malloc(sizeof(Node)*no_of_nodes);
bool *h_graph_mask = (bool*) malloc(sizeof(bool)*no_of_nodes);
bool *h_updating_graph_mask = (bool*) malloc(sizeof(bool)*no_of_nodes);
bool *h_graph_visited = (bool*) malloc(sizeof(bool)*no_of_nodes);

int start, edgeno;
// initialize the memory
for( unsigned int i = 0; i < no_of_nodes; i++)
{
    fscanf(fp,"%d %d",&start,&edgeno);
    h_graph_nodes[i].starting = start;
    h_graph_nodes[i].no_of_edges = edgeno;
    h_graph_mask[i]=false;
    h_updating_graph_mask[i]=false;
    h_graph_visited[i]=false;
}

//read the source node from the file
fscanf(fp,"%d",&source);
//
source=0;

//set the source node as true in the mask
h_graph_mask[source]=true;
h_graph_visited[source]=true;

fscanf(fp,"%d",&edge_list_size);

int id,cost;
int* h_graph_edges = (int*) malloc(sizeof(int)*edge_list_size);
for(int i=0; i < edge_list_size ; i++)
{
    fscanf(fp,"%d",&id);
    fscanf(fp,"%d",&cost);
    h_graph_edges[i] = id;
}

if(fp)
    fclose(fp);

printf("Read File\n");

//Copy the Node list to device memory
Node* d_graph_nodes;
cudaMalloc( (void**) &d_graph_nodes, sizeof(Node)*no_of_nodes) ;

```

```

    cudaMemcpy( d_graph_nodes, h_graph_nodes, sizeof(Node)*no_of_nodes,
cudaMemcpyHostToDevice) ;

    //Copy the Edge List to device Memory
    int* d_graph_edges;
    cudaMalloc( (void**) &d_graph_edges, sizeof(int)*edge_list_size) ;
    cudaMemcpy( d_graph_edges, h_graph_edges, sizeof(int)*edge_list_size,
cudaMemcpyHostToDevice) ;

    //Copy the Mask to device memory
    bool* d_graph_mask;
    cudaMalloc( (void**) &d_graph_mask, sizeof(bool)*no_of_nodes) ;
    cudaMemcpy( d_graph_mask, h_graph_mask, sizeof(bool)*no_of_nodes,
cudaMemcpyHostToDevice) ;

    bool* d_updating_graph_mask;
    cudaMalloc( (void**) &d_updating_graph_mask, sizeof(bool)*no_of_nodes) ;
    cudaMemcpy( d_updating_graph_mask, h_updating_graph_mask,
sizeof(bool)*no_of_nodes, cudaMemcpyHostToDevice) ;

    //Copy the Visited nodes array to device memory
    bool* d_graph_visited;
    cudaMalloc( (void**) &d_graph_visited, sizeof(bool)*no_of_nodes) ;
    cudaMemcpy( d_graph_visited, h_graph_visited, sizeof(bool)*no_of_nodes,
cudaMemcpyHostToDevice) ;

    // allocate mem for the result on host side
    int* h_cost = (int*) malloc( sizeof(int)*no_of_nodes);
    for(int i=0;i<no_of_nodes;i++)
        h_cost[i]=-1;
    h_cost[source]=0;

    // allocate device memory for result
    int* d_cost;
    cudaMalloc( (void**) &d_cost, sizeof(int)*no_of_nodes);
    cudaMemcpy( d_cost, h_cost, sizeof(int)*no_of_nodes,
cudaMemcpyHostToDevice) ;

    //make a bool to check if the execution is over
    bool *d_over;
    cudaMalloc( (void**) &d_over, sizeof(bool));

    printf("Copied Everything to GPU memory\n");

    // setup execution parameters
    dim3 grid( num_of_blocks, 1, 1);
    dim3 threads( num_of_threads_per_block, 1, 1);

```

```

int k=0;
printf("Start traversing the tree\n");
bool stop;
//Call the Kernel untill all the elements of Frontier are not false
do
{
    //if no thread changes this value then the loop stops
    stop=false;
    cudaMemcpy( d_over, &stop, sizeof(bool), cudaMemcpyHostToDevice) ;
    Kernel<<< grid, threads, 0 >>>( d_graph_nodes, d_graph_edges,
d_graph_mask, d_updating_graph_mask, d_graph_visited, d_cost, no_of_nodes);
    // check if kernel execution generated and error

    Kernel2<<< grid, threads, 0 >>>( d_graph_mask,
d_updating_graph_mask, d_graph_visited, d_over, no_of_nodes);
    // check if kernel execution generated and error

    cudaMemcpy( &stop, d_over, sizeof(bool), cudaMemcpyDeviceToHost) ;
    k++;
}
while(stop);

printf("Kernel Executed %d times\n",k);

// copy result from device to host
cudaMemcpy( h_cost, d_cost, sizeof(int)*no_of_nodes,
cudaMemcpyDeviceToHost) ;

//Store the result into a file
FILE *fpo = fopen("result.txt","w");
for(int i=0;i<no_of_nodes;i++)
    fprintf(fpo,"%d cost:%d\n",i,h_cost[i]);
fclose(fpo);
printf("Result stored in result.txt\n");

// cleanup memory
free( h_graph_nodes);
free( h_graph_edges);
free( h_graph_mask);
free( h_updating_graph_mask);
free( h_graph_visited);
free( h_cost);
cudaFree(d_graph_nodes);
cudaFree(d_graph_edges);

```

```
    cudaFree(d_graph_mask);
    cudaFree(d_updating_graph_mask);
    cudaFree(d_graph_visited);
    cudaFree(d_cost);
}

int main( int argc, char** argv)
{
    no_of_nodes=0;
    edge_list_size=0;
    BFS();
}
```

C.2. Parallel st-CON

```

/*****
Implementing s-t Connectivity algorithm on CUDA
*****/
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <math.h>
#include <cuda.h>

#define MAX_THREADS_PER_BLOCK 512

int no_of_nodes;
int edge_list_size;
FILE *fp;

//Structure to hold a node information
struct Node
{
    int starting;
    int no_of_edges;
};

__device__ int rf=0;
__device__ int gf=0;
__device__ bool d_stop=false;

__global__ void
Kernel( Node* g_graph_nodes, int* g_graph_edges, bool* g_graph_mask,
        bool* g_updating_graph_mask, bool *g_Red_graph_visited,
        bool *g_Green_graph_visited, bool *g_Red_updating_graph_visited,
        bool *g_Green_updating_graph_visited, int* g_cost, int no_of_nodes,
        bool* g_over)
{
    int tid = blockIdx.x*blockDim.x + threadIdx.x;
    if( tid<no_of_nodes && g_graph_mask[tid])
    {
        g_graph_mask[tid]=false;
        for(int i=g_graph_nodes[tid].starting;
            i<(g_graph_nodes[tid].no_of_edges + g_graph_nodes[tid].starting); i++)
        {
            int nid = g_graph_edges[i];
            if(g_Red_graph_visited[nid] || g_Green_graph_visited[nid])
            {
                if(g_Red_graph_visited[tid] &&
                    g_Green_graph_visited[nid])
            }
        }
    }
}

```



```

        {
            rf = g_cost[tid]+1;
            *g_over = true;
            d_stop=true;
        }
        if(g_Green_graph_visited[tid] &&
            g_Red_graph_visited[nid])
        {
            *g_over=true;
            d_stop=true;
        }
    }
    else
    {
        if(g_Green_graph_visited[tid])
            g_Green_updating_graph_visited[nid] = true;
        if(g_Red_graph_visited[tid])
            g_Red_updating_graph_visited[nid] = true;

        g_updating_graph_mask[nid]=true;
        g_cost[nid] = g_cost[tid]+1;
    }
}

}

}

__global__ void
Kernel2( bool* g_graph_mask, bool * g_updating_graph_mask, bool *
    g_Red_graph_visited, bool * g_Green_graph_visited, bool *
    g_Red_updating_graph_visited, bool * g_Green_updating_graph_visited, bool
    *g_over, int no_of_nodes,int* g_cost)
{
    int tid = blockIdx.x*blockDim.x + threadIdx.x;
    if( tid<no_of_nodes && g_updating_graph_mask[tid] && !d_stop)
    {
        //printf("\n%dT",tid);
        g_graph_mask[tid]=true;
        if(g_Red_updating_graph_visited[tid])
        {
            g_Red_graph_visited[tid] = true;
            rf = g_cost[tid];
            g_Red_updating_graph_visited[tid] = false;
        }
        if(g_Green_updating_graph_visited[tid])
        {
            g_Green_graph_visited[tid] = true;

```

```

        gf = g_cost[tid];
        g_Green_updating_graph_visited[tid] = false;
    }
    g_updating_graph_mask[tid]=false;

    if(g_Green_graph_visited[tid] && g_Red_graph_visited[tid])
        *g_over=true;
}

}

__global__
void dummy(int *len){
    //printf("R%dG%d",rf,gf);
    *len = rf+gf;
}

////////////////////////////////////
//Apply BFS on a Graph using CUDA
////////////////////////////////////
void stCON() {
    char *input_f;
    input_f = (char*)malloc(sizeof(char)*100);
    printf("File path:");
    scanf("%s",input_f);

    printf("Reading File\n");
    //Read in Graph from a file
    fp = fopen(input_f,"r");
    if(!fp)
    {
        printf("Error Reading graph file\n");
        return;
    }

    int source = 0;
    int terminal = 0; //terminal (Appended)
    //int h_R_length = 0; //Rf (Appended)
    //int h_G_length = 0; //Gf (Appended)

    fscanf(fp,"%d",&no_of_nodes);

    int num_of_blocks = 1;
    int num_of_threads_per_block = no_of_nodes;

    //Make execution Parameters according to the number of nodes
    //Distribute threads across multiple Blocks if necessary
    if(no_of_nodes>MAX_THREADS_PER_BLOCK)
    {

```

```

        num_of_blocks =
(int)ceil(no_of_nodes/(double)MAX_THREADS_PER_BLOCK);
        num_of_threads_per_block = MAX_THREADS_PER_BLOCK;
    }

    // allocate host memory
    Node* h_graph_nodes = (Node*) malloc(sizeof(Node)*no_of_nodes);
    bool *h_graph_mask = (bool*) malloc(sizeof(bool)*no_of_nodes);
    //Fa
    bool *h_updating_graph_mask = (bool*) malloc(sizeof(bool)*no_of_nodes);
    //Fva
    bool *h_Red_graph_visited = (bool*) malloc(sizeof(bool)*no_of_nodes);
    //Ra (Appended)
    bool *h_Green_graph_visited = (bool*) malloc(sizeof(bool)*no_of_nodes);
    //Ga (Appended)
    bool *h_Red_updating_graph_visited = (bool*)
malloc(sizeof(bool)*no_of_nodes); //Rva (Appended)
    bool *h_Green_updating_graph_visited = (bool*)
malloc(sizeof(bool)*no_of_nodes); //Gva (Appended)

    int start, edgeno;
    // initialize the memory
    for( unsigned int i = 0; i < no_of_nodes; i++)
    {
        fscanf(fp,"%d %d",&start,&edgeno);
        h_graph_nodes[i].starting = start;
        h_graph_nodes[i].no_of_edges = edgeno;
        h_graph_mask[i]=false;
        h_updating_graph_mask[i]=false;
        h_Red_graph_visited[i]=false; //Ra=false (Appended)
        h_Green_graph_visited[i]=false; //Ga=false (Appended)
        h_Red_updating_graph_visited[i]=false; //Rva=false
(Appended)
        h_Green_updating_graph_visited[i]=false; //Gva=false (Appended)
    }

    //read the source node from the file
    fscanf(fp,"%d",&source);
//
    source=0;
    fscanf(fp,"%d",&terminal); //take input terminal from file (Appended)

    //set the source node as true in the mask
    h_graph_mask[source]=true;
    h_graph_mask[terminal]=true; //Added line
    h_Red_graph_visited[source]=true; //Appended line (Ra)
    h_Green_graph_visited[terminal]=true; //Appended line (Ga)

    fscanf(fp,"%d",&edge_list_size);

```

```

int id,cost;
int* h_graph_edges = (int*) malloc(sizeof(int)*edge_list_size);
for(int i=0; i < edge_list_size ; i++)
{
    fscanf(fp,"%d",&id);
    fscanf(fp,"%d",&cost); //needed because of the format of the file
    h_graph_edges[i] = id;
}

if(fp)
    fclose(fp);

printf("Read File\n");

//Copy the Node list to device memory
Node* d_graph_nodes;
cudaMalloc( (void**) &d_graph_nodes, sizeof(Node)*no_of_nodes) ;
cudaMemcpy( d_graph_nodes, h_graph_nodes, sizeof(Node)*no_of_nodes,
cudaMemcpyHostToDevice) ;

//Copy the Edge List to device Memory
int* d_graph_edges;
cudaMalloc( (void**) &d_graph_edges, sizeof(int)*edge_list_size) ;
cudaMemcpy( d_graph_edges, h_graph_edges, sizeof(int)*edge_list_size,
cudaMemcpyHostToDevice) ;

//Copy the Mask to device memory
bool* d_graph_mask;
cudaMalloc( (void**) &d_graph_mask, sizeof(bool)*no_of_nodes) ;
cudaMemcpy( d_graph_mask, h_graph_mask, sizeof(bool)*no_of_nodes,
cudaMemcpyHostToDevice) ;

bool* d_updating_graph_mask;
cudaMalloc( (void**) &d_updating_graph_mask, sizeof(bool)*no_of_nodes) ;
cudaMemcpy( d_updating_graph_mask, h_updating_graph_mask,
sizeof(bool)*no_of_nodes, cudaMemcpyHostToDevice) ;

//Copy the Visited nodes array to device memory (Appended)
bool* d_Red_graph_visited;
cudaMalloc( (void**) &d_Red_graph_visited, sizeof(bool)*no_of_nodes) ;
cudaMemcpy( d_Red_graph_visited, h_Red_graph_visited,
sizeof(bool)*no_of_nodes, cudaMemcpyHostToDevice) ;

//Copy the Visited nodes array to device memory (Appended)
bool* d_Green_graph_visited;
cudaMalloc( (void**) &d_Green_graph_visited, sizeof(bool)*no_of_nodes) ;

```

```

    cudaMemcpy( d_Green_graph_visited, h_Green_graph_visited,
sizeof(bool)*no_of_nodes, cudaMemcpyHostToDevice) ;

    //Copy the Visited nodes array to device memory (Appended)
    bool* d_Red_updating_graph_visited;
    cudaMalloc( (void**) &d_Red_updating_graph_visited,
sizeof(bool)*no_of_nodes) ;
    cudaMemcpy( d_Red_updating_graph_visited, h_Red_updating_graph_visited,
sizeof(bool)*no_of_nodes, cudaMemcpyHostToDevice) ;

    //Copy the Visited nodes array to device memory (Appended)
    bool* d_Green_updating_graph_visited;
    cudaMalloc( (void**) &d_Green_updating_graph_visited,
sizeof(bool)*no_of_nodes) ;
    cudaMemcpy( d_Green_updating_graph_visited,
h_Green_updating_graph_visited, sizeof(bool)*no_of_nodes,
cudaMemcpyHostToDevice) ;

    // allocate mem for the result on host side
    int* h_cost = (int*) malloc( sizeof(int)*no_of_nodes);
    for(int i=0;i<no_of_nodes;i++)
        h_cost[i]=-1;
    h_cost[source]=0;
    h_cost[terminal]=0;          //(Appended)

    // allocate device memory for result
    int* d_cost;
    cudaMalloc( (void**) &d_cost, sizeof(int)*no_of_nodes);
    cudaMemcpy( d_cost, h_cost, sizeof(int)*no_of_nodes,
cudaMemcpyHostToDevice) ;

    //make a bool to check if the execution is over
    bool *d_over;
    cudaMalloc( (void**) &d_over, sizeof(bool));

    printf("Copied Everything to GPU memory\n");

    // setup execution parameters
    dim3 grid( num_of_blocks, 1, 1);
    dim3 threads( num_of_threads_per_block, 1, 1);

    int k=0;
    printf("Start traversing the tree\n");
    bool stop=false;
    cudaMemcpy( d_over, &stop, sizeof(bool), cudaMemcpyHostToDevice) ;
    //Call the Kernel untill all the elements of Frontier are not false
    do
    {

```

```

    Kernel<<< grid, threads, 0 >>>(d_graph_nodes, d_graph_edges,
        d_graph_mask, d_updating_graph_mask,
        d_Red_graph_visited, d_Green_graph_visited,
        d_Red_updating_graph_visited, d_Green_updating_graph_visited,
        d_cost, no_of_nodes, d_over);
    // check if kernel execution generated and error

    Kernel2<<< grid, threads, 0 >>>(d_graph_mask, d_updating_graph_mask,
        d_Red_graph_visited, d_Green_graph_visited,
        d_Red_updating_graph_visited, d_Green_updating_graph_visited,
        d_over, no_of_nodes, d_cost);
    // check if kernel execution generated and error

    cudaMemcpy( &stop, d_over, sizeof(bool), cudaMemcpyDeviceToHost) ;
    k++;
}
while(!stop);

printf("Kernel Executed %d times\n",k);

// copy result from device to host
cudaMemcpy( h_cost, d_cost, sizeof(int)*no_of_nodes,
cudaMemcpyDeviceToHost) ;

int *d_len,len;
cudaMalloc( (void**) &d_len, sizeof(int));
dummy<<<1,1>>>(d_len);
cudaMemcpy( &len, d_len, sizeof(int), cudaMemcpyDeviceToHost);
printf("\nlength == %d",len);

//Store the result into a file
FILE *fpo = fopen("result.txt","w");
for(int i=0;i<no_of_nodes;i++)
    fprintf(fpo,"%d cost:%d\n",i,h_cost[i]);
fclose(fpo);
printf("Result stored in result.txt\n");

// cleanup memory
free( h_graph_nodes);
free( h_graph_edges);
free( h_graph_mask);

```

```

    free( h_updating_graph_mask);
    free( h_Red_graph_visited);          //(Appended)
    free( h_Green_graph_visited); //(Appended)
    free( h_Red_updating_graph_visited);  //(Appended)
    free( h_Green_updating_graph_visited); //(Appended)
    free( h_cost);
    cudaFree(d_graph_nodes);
    cudaFree(d_graph_edges);
    cudaFree(d_graph_mask);
    cudaFree(d_updating_graph_mask);
    cudaFree(d_Red_graph_visited);        //(Appended)
    cudaFree(d_Green_graph_visited);      //(Appended)
    cudaFree(d_Red_updating_graph_visited); //(Appended)
    cudaFree(d_Green_updating_graph_visited); //(Appended)

//    cudaFree(d_graph_visited);
    cudaFree(d_cost);
    cudaFree(d_len);
}

int main( int argc, char** argv) {
    no_of_nodes=0;
    edge_list_size=0;
    stCON();
}

```

C.3. Parallel SSSP

```
/*
Implementing Single Source Shortest Path on CUDA
*/
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <math.h>
#include <assert.h>

#define MAX_COST 10000000

unsigned int no_of_nodes;
unsigned int edge_list_size;
FILE *fp;

__global__ void
DijkstraKernel1(unsigned int* g_graph_nodes, unsigned int* g_graph_edges,
    unsigned int* g_graph_weights, unsigned int* g_graph_updating_cost, bool*
    g_graph_mask, unsigned int* g_cost , unsigned int no_of_nodes,
    unsigned int edge_list_size)
{
    unsigned int tid = blockIdx.x*blockDim.x + threadIdx.x;
    unsigned int i,id;
    unsigned int end = edge_list_size;
    if(tid<no_of_nodes && g_graph_mask[tid])
    {
        if(tid < no_of_nodes-1)
            end = g_graph_nodes[tid+1];
        for(i = g_graph_nodes[tid]; i< end; i++)
        {
            id = g_graph_edges[i];
            atomicMin(&g_graph_updating_cost[id],
                g_cost[tid]+g_graph_weights[i]);
        }
        g_graph_mask[tid]=false;
    }
}

__global__ void
DijkstraKernel2(unsigned int* g_graph_nodes, unsigned int* g_graph_edges,
    unsigned int* g_graph_weights, unsigned int* g_graph_updating_cost, bool*
    g_graph_mask, unsigned int* g_cost , bool *d_finished, unsigned int
    no_of_nodes, unsigned int edge_list_size)
{
    unsigned int tid = blockIdx.x*blockDim.x + threadIdx.x;
```



```

        if(tid<no_of_nodes && g_cost[tid] > g_graph_updating_cost[tid])
        {
            g_cost[tid] = g_graph_updating_cost[tid];
            g_graph_mask[tid] = true;
            *d_finished = true;
        }
        if(tid<no_of_nodes)
            g_graph_updating_cost[tid] = g_cost[tid];
    }

    //////////////////////////////////////
    //Apply Shortest Path on a Graph using CUDA
    //////////////////////////////////////
    void Dijkstra()
    {
        printf("Reading File\n");
        fp = fopen("g.txt","r");
        if(!fp)
        {
            printf("Error Reading graph file\n");
            return;
        }

        unsigned int source = 0;

        fscanf(fp,"%d",&no_of_nodes);
        printf("No of Nodes: %d\n",no_of_nodes);

        cudaDeviceProp dev;
        cudaGetDeviceProperties(&dev,0);
        //printf("thread = %d",dev.maxThreadsPerBlock);

        unsigned int MAX_THREADS_PER_BLOCK;
        unsigned int num_of_blocks = 1;
        unsigned int num_of_threads_per_block = no_of_nodes;
        MAX_THREADS_PER_BLOCK = dev.maxThreadsPerBlock;

        //Make execution Parameters according to the number of nodes
        //Distribute threads across multiple Blocks if necessary
        if(no_of_nodes>MAX_THREADS_PER_BLOCK)
        {
            num_of_blocks = (unsigned int)ceil(no_of_nodes
                                                /(double)MAX_THREADS_PER_BLOCK);
            num_of_threads_per_block = MAX_THREADS_PER_BLOCK;
        }

        // allocate host memory

```

```

    unsigned int* h_graph_nodes = (unsigned int*) malloc(sizeof(unsigned
int)*no_of_nodes);
    bool *h_graph_mask = (bool*) malloc(sizeof(bool)*no_of_nodes);
    unsigned int *h_graph_updating_cost = (unsigned int*) malloc(sizeof(unsigned
int)*no_of_nodes);

    unsigned int start, edgeno;
    // initialize the memory
    unsigned int no=0;
    for( unsigned int i = 0; i < no_of_nodes; i++)
    {
        fscanf(fp,"%d %d",&start,&edgeno);
        if(edgeno>100)
            no++;
        h_graph_nodes[i] = start;
        h_graph_updating_cost[i] = MAX_COST;
        h_graph_mask[i]=false;
    }

    //read the source unsigned int from the file
    fscanf(fp,"%d",&source);
    printf("Source %d\n",source);

    //set the source unsigned int as true in the mask
    h_graph_mask[source]=true;
    //h_graph_counter[source]=0;

    fscanf(fp,"%d",&edge_list_size);

    unsigned int id;
    unsigned int* h_graph_edges = (unsigned int*) malloc(sizeof(unsigned
int)*edge_list_size);
    unsigned int* h_graph_weights = ( unsigned int*) malloc(sizeof( unsigned
int)*edge_list_size);

    unsigned int i;
    for(i=0; i < edge_list_size ; i++)
    {
        fscanf(fp,"%d",&id);
        h_graph_edges[i] = id;
        fscanf(fp,"%d",&id);
        h_graph_weights[i] = id;
    }
    if(fp)
        fclose(fp);

    printf("Read File\n");

```

```

    //Copy the unsigned int list to device memory
    unsigned int* d_graph_nodes;
    cudaMalloc( (void**) &d_graph_nodes, sizeof(unsigned int)*no_of_nodes);
    cudaMemcpy( d_graph_nodes, h_graph_nodes, sizeof(unsigned int)*no_of_nodes,
cudaMemcpyHostToDevice);

    //Copy the Edge List to device Memory
    unsigned int* d_graph_edges;
    cudaMalloc( (void**) &d_graph_edges, sizeof(unsigned int)*edge_list_size);
    cudaMemcpy( d_graph_edges, h_graph_edges, sizeof(unsigned
int)*edge_list_size, cudaMemcpyHostToDevice);

    unsigned int* d_graph_weights;
    cudaMalloc( (void**) &d_graph_weights, sizeof( unsigned
int)*edge_list_size);
    cudaMemcpy( d_graph_weights, h_graph_weights, sizeof( unsigned
int)*edge_list_size, cudaMemcpyHostToDevice);

    //Copy the Mask to device memory
    bool* d_graph_mask;
    cudaMalloc( (void**) &d_graph_mask, sizeof(bool)*no_of_nodes);
    cudaMemcpy( d_graph_mask, h_graph_mask, sizeof(bool)*no_of_nodes,
cudaMemcpyHostToDevice);

    // allocate mem for the result on host side
    unsigned int* h_cost = (unsigned int*) malloc( sizeof(unsigned
int)*no_of_nodes);
    for(unsigned int i=0;i<no_of_nodes;i++)
        h_graph_updating_cost[i] = h_cost[i] = MAX_COST;
    h_cost[source]=0;
    // allocate device memory for result
    unsigned int* d_cost;
    cudaMalloc( (void**) &d_cost, sizeof(unsigned int)*no_of_nodes);
    cudaMemcpy( d_cost, h_cost, sizeof(unsigned int)*no_of_nodes,
cudaMemcpyHostToDevice);

    unsigned int* d_graph_updating_cost;
    cudaMalloc( (void**) &d_graph_updating_cost, sizeof(unsigned
int)*no_of_nodes);
    cudaMemcpy( d_graph_updating_cost, h_graph_updating_cost, sizeof(unsigned
int)*no_of_nodes, cudaMemcpyHostToDevice);

    //make a bool to check if the execution is over

    bool *d_finished;
    bool finished;
    cudaMalloc( (void**) &d_finished, sizeof(bool));

```

```

// setup execution parameters
dim3 grid( num_of_blocks, 1, 1);
dim3 threads( num_of_threads_per_block, 1, 1);

unsigned int k=0;

do
{
    DijkstraKernel1<<< grid, threads, 0 >>>( d_graph_nodes,
        d_graph_edges, d_graph_weights, d_graph_updating_cost,
        d_graph_mask, d_cost, no_of_nodes, edge_list_size);
    cudaDeviceSynchronize();
    k++;
    finished=false;
    cudaMemcpy( d_finished, &finished, sizeof(bool),
cudaMemcpyHostToDevice);
    DijkstraKernel2<<< grid, threads, 0 >>>( d_graph_nodes,
        d_graph_edges, d_graph_weights, d_graph_updating_cost,
        d_graph_mask, d_cost, d_finished, no_of_nodes, edge_list_size);
    cudaDeviceSynchronize();
    cudaMemcpy( &finished, d_finished, sizeof(bool),
cudaMemcpyDeviceToHost);
}
while(finished);

// copy result from device to host
cudaMemcpy( h_cost, d_cost, sizeof(unsigned int)*no_of_nodes,
cudaMemcpyDeviceToHost);
cudaDeviceSynchronize();

//Store the result unsigned into a file
FILE *fpo = fopen("result.txt","w");
for(unsigned int i=0;i<no_of_nodes;i++)
fprintf(fpo,"%d) cost:%d\n",i,h_cost[i]);
fclose(fpo);
printf("Result stored in result.txt\n");

// cleanup memory
free( h_graph_nodes);
free( h_graph_edges);
free( h_graph_mask);
free( h_graph_weights);
free( h_graph_updating_cost);
free( h_cost);
cudaFree(d_graph_nodes);
cudaFree(d_graph_edges);

```

```
    cudaFree(d_graph_mask);
    cudaFree(d_graph_weights);
    cudaFree(d_graph_updating_cost);
    cudaFree(d_cost);
    cudaFree(d_finished);
}

int main( int argc, char** argv)
{
    no_of_nodes=0;
    edge_list_size=0;
    Dijkstra();
}
```

C.4. Parallel 8-Puzzle Solver

```
/*
Implementing parallel 8-puzzle solver on CUDA
*/
#include <stdio.h>
#include <assert.h>

#define LEN 9
#define SIZE 362880 //9!

#define CUDA_CHECK_RETURN(value) { \
    cudaError_t _m_cudaStat = value; \
    if (_m_cudaStat != cudaSuccess) { \
        fprintf(stderr, "Error %s at line %d in file %s\n", \
            cudaGetErrorString(_m_cudaStat), __LINE__, __FILE__); \
        exit(1); \
    } \
}

enum Move{UP, DOWN, LEFT, RIGHT};

char s[LEN+1], g[LEN+1];
bool *frontier, *Ufrontier;
int *visited;

__device__ char *start, *goal;
__device__ int sIndex, gIndex;
__device__ char *state, *charH, *charT;

__device__
int fact(int n){ //computes factorial
    int x=1;
    for (int i = 1; i <= n; i++) {
        x *= i;
    }
    return x;
}

//Implementing minimal perfect hashing
//This will generate the index/rank of a given permutation
__device__
int getHash(char *n,int tid){ //implements perfect hashing
    int h=0;
    char *num;
    num = &charH[tid*9];
```

```

    for (int i = 0; i < LEN; ++i) {
        num[i] = n[i] - '0';
    }

    int f=LEN;
    for (int i = 0; i < LEN ; ++i) {
        f--;
        if(num[i] > 0)
            h += num[i]*fact(f);
        for (int j = i+1; j < LEN; ++j) {
            if(num[j] > num[i]){
                num[j]--;
            }
        }
    }
    return h;
}

```

__global__

```

void init(char *s, char *g, int *visited, bool *frontier, bool *Ufrontier, char
        *h, char *t, char *st){
    start=s;
    goal=g;

    charH = h
    if(charH==NULL){
        printf("charH NULL");
        assert(0);
    }
    charT = t;
    if(charT==NULL){
        printf("charT NULL");
        assert(0);
    }
    state = st;
    if(state==NULL){
        printf("state NULL");
        assert(0);
    }

    sIndex=getHash(s,0);
    int offset = sIndex*9;
    for (int i = 0; i < LEN; ++i) {
        state[offset+i] = s[i];
    }

    gIndex=getHash(g,0);

```

```

        if(visited==NULL){
            printf("visited NULL");
            assert(0);
        }
        if(frontier==NULL){
            printf("frontier NULL");
            assert(0);
        }
        if(Ufrontier==NULL){
            printf("Ufrontier NULL");
            assert(0);
        }
    }
}

__global__
void clean(int *visited, bool *frontier, bool *Ufrontier){
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    if(tid < SIZE){
        visited[tid] = -1;
        frontier[tid] = false;
        Ufrontier[tid] = false;
    }
}

__device__
char* swap(char *c, int pos, int p, int offset){
    char *a;
    a = charT+offset;//(char*)malloc(sizeof(char)*LEN);
    if(a==NULL){
        printf("a NULL");
        assert(0);
    }
    for (int i = 0; i < LEN; ++i) {
        a[i] = c[offset+i];
    }
    int x=pos;
    int y=pos+p;
    char tmp = a[x];
    a[x] = a[y];
    a[y] = tmp;
    return a;
}

__device__
char* move(char *s, int pos, Move m, int tid){
    int offset = tid*9;
    int i,j;

```



```

i=pos/3;
j=pos%3;
switch(m){
    case UP:
        if(i==0)
            return NULL;
        return swap(s, pos, -3, offset);
    case DOWN:
        if(i==2)
            return NULL;
        return swap(s, pos, 3, offset);
    case LEFT:
        if(j==0)
            return NULL;
        return swap(s, pos, -1, offset);
    case RIGHT:
        if(j==2)
            return NULL;
        return swap(s, pos, 1, offset);

    default: return NULL;
}
}

__global__
void compute(int *visited, bool *frontier, bool *Ufrontier, bool *fin){
    char *adj;
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    int index, pos;
    int offset = tid*9;
    if(tid < SIZE && frontier[tid]){
        frontier[tid] = false;
        if(tid==gIndex)
            *fin = true;
        for (int i = 0; i < LEN; ++i) {
            if(state[offset+i] == '0')
                pos = i;
        }
        for (int i = UP; i <= RIGHT; ++i) {
            adj=move(state, pos, (Move)i, tid);
            if(adj == NULL){
                continue;
            }
            index = getHash(adj,tid);
            offset = index*9;
            if(visited[index] < 0){
                Ufrontier[index] = true;
            }
        }
    }
}

```

```

        visited[index] = tid;
        for (int i = 0; i < LEN; ++i) {
            state[offset+i] = adj[i];
        }
    }
}

__global__
void save(bool *frontier, bool *Ufrontier, bool *fin){
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    if(tid < SIZE && Ufrontier[tid]){
        Ufrontier[tid] = false;
        frontier[tid] = true;
    }
}

__global__
void dummy(bool *frontier, bool *fin){
    frontier[sIndex] = true;
}

int main(int argc, char **argv) {
    printf("Enter start: ");
    scanf("%s",s);

    printf("Enter goal: ");
    scanf("%s",g);

    char *start, *goal;
    cudaMalloc((void**)&start,sizeof(char)*10);
    cudaMalloc((void**)&goal,sizeof(char)*10);

    cudaMemcpy(start,&s,sizeof(char)*10, cudaMemcpyHostToDevice);
    cudaMemcpy(goal,&g,sizeof(char)*10, cudaMemcpyHostToDevice);

    cudaMalloc((void**)&visited,sizeof(int)*SIZE);
    cudaMalloc((void**)&frontier,sizeof(bool)*SIZE);
    cudaMalloc((void**)&Ufrontier,sizeof(bool)*SIZE);
    char *h, *t, *st;
    cudaMalloc((void**)&h,sizeof(char)*SIZE*9);
    cudaMalloc((void**)&t,sizeof(char)*SIZE*9);
    cudaMalloc((void**)&st,sizeof(char)*SIZE*9);

    int threads = 504;
    int blocks = SIZE/threads;

```

```

init<<<1,1>>>(start, goal, visited, frontier, Ufrontier, h, t, st);
CUDA_CHECK_RETURN(cudaDeviceSynchronize());
clean<<<blocks, threads>>>(visited, frontier, Ufrontier);
CUDA_CHECK_RETURN(cudaDeviceSynchronize());

bool fin = false, *dfin;
cudaMalloc((void**)&dfin, sizeof(bool));
cudaMemcpy(dfin, &fin, sizeof(bool), cudaMemcpyHostToDevice);
dummy<<<1,1>>>(frontier, dfin);
CUDA_CHECK_RETURN(cudaDeviceSynchronize());
int k=0;
while(!fin){
    compute<<<blocks, threads>>>(visited, frontier, Ufrontier, dfin);
    CUDA_CHECK_RETURN(cudaDeviceSynchronize());
    save<<<blocks, threads>>>(frontier, Ufrontier, dfin);
    cudaMemcpy(&fin, dfin, sizeof(bool), cudaMemcpyDeviceToHost);
    CUDA_CHECK_RETURN(cudaDeviceSynchronize());
    k++;
}
printf("\n%d", k-1);

//free allocated memories

}

```

C.5. Complete 8-Puzzle Solver

```

/*****
Implementing parallel 8-puzzle solver on CUDA
*****/
#include <ctime>
#include <iostream>
#include <stdio.h>
#include <assert.h>

#define CUDA_CHECK_RETURN(value) {                                     \
    cudaError_t _m_cudaStat = value;                                  \
    if (_m_cudaStat != cudaSuccess) {                                  \
        fprintf(stderr, "Error %s at line %d in file %s\n",          \
            cudaGetErrorString(_m_cudaStat), __LINE__, __FILE__); \
        exit(1);                                                      \
    } }

#define LEN 9
#define mySIZE 362880 //9!

using namespace std;

class Queue{
private:
    int front,rear;
    int qarr[30000]; //a heuristic value is taken
    int s;
public:
    __device__
    void init(){
        if(&qarr == NULL){
            printf("Queue failed\n");
            assert(0);
            return;
        }
        front = 0;
        rear = 0;
        s=0;
    }

    __device__
    void enqueue(int x){
        if(s==30000){
            printf("Q full");
            assert(0);
        }
    }
}

```

```

        qarr[rear] = x;
        rear++;
        if(rear == 30000)
            rear = 0;
        s++;
    }

    __device__
    int deQueue(){
        if(isEmpty())//{
            return -1; //empty
        }
        s--;
        int x = qarr[front];
        front++;
        if(front == 30000)
            front = 0;
        return x;
    }

    __device__
    bool isEmpty(){
        if(s == 0)
            return true;
        return false;
    }

    __device__
    int size(){
        return s;
    }
};

__device__
int fact(int n){
    int x=1;
    for (int i = 1; i <= n; i++) {
        x *= i;
    }
    return x;
}

__device__ int factof[9];

__global__
void storeFact(){
    for (int i = 0; i < LEN; ++i) {
        factof[i] = fact(i);
    }
}

```

```

////////////////////////////////////
//Implementing minimal perfect hashing
//This will generate the index/rank of a given permutation
////////////////////////////////////
__device__
int getHash(int n, char *numH){
    int h=0;
    for (int i = LEN-1; i >= 0; --i) {
        numH[i] = n%10;
        n = n/10;
    }
    int f=LEN;
    for (int i = 0; i < LEN ; ++i) {
        f--;
        if(numH[i] > 0)
            h += numH[i]*factof[f];
        for (int j = i+1; j < LEN; ++j) {
            if(numH[j] > numH[i]){
                numH[j]--;
            }
        }
    }
    return h;
}

__device__ int state[mySIZE];
__device__ int k = 0;

__device__ __noinline__
void generate(int x[], int l){
    if (l == LEN) {
        int num = 0;
        for (int i = 0; i < LEN; i++) {
            num = num * 10 + x[i];
        }
        state[k++] = num;
        return;
    }
    int j;
    for (int i = 0; i < LEN; i++) {
        for (j = 0; j < l; j++) {
            if(x[j] == i)
                break; //next i
        }
        if(j==l){ //fully iterated, i not in x[]
            x[l] = i;
            generate(x,l+1);
        }
    }
}

```

```

    }
}

__global__ __noinline__
void populate()
{
    int x[LEN] = { 0, 1, 2, 3, 4, 5, 6, 7, 8 };
    generate(x,0);
}

enum Move{UP, DOWN, LEFT, RIGHT};

__device__ int cost[mySIZE];
__device__ int count = 0;
__device__ int gIndex;

__device__
void dump(char str[]){
    for (int i = 0; i < 9; ++i) {
        printf("%d",str[i]);
    }
    printf("\n");
}

__device__
char* swap(char str[],int pos,int p){
    int x=pos;
    int y=pos+p;
    char tmp = str[x];
    str[x] = str[y];
    str[y] = tmp;
    pos = y;
    return str;
}

__global__
void setGoal(int n){
    char tmp[LEN];
    gIndex = getHash(n,tmp);
}

__device__
int move(int state,Move m,char *nextM){
    int pos,i,j;
    for (int k = LEN-1; k >= 0; --k) {
        nextM[k] = state%10;
        state = state/10;
    }
}

```

```

        if(nextM[k] == 0)
            pos = k;
    }
    i=pos/3;
    j=pos%3;
    switch(m){
        case UP:
            if(i==0)
                return -1;
            swap(nextM,pos,-3);
            break;
        case DOWN:
            if(i==2)
                return -1;
            swap(nextM,pos,3);
            break;
        case LEFT:
            if(j==0)
                return -1;
            swap(nextM,pos,-1);
            break;
        case RIGHT:
            if(j==2)
                return -1;
            swap(nextM,pos,1);
            break;
        default: return -1;
    }
    int x=0;
    for (int i = 0; i < LEN; ++i) {
        x = x*10 + nextM[i];
    }
    return x;
}

```

```

class eightPuzzle{
private:
    Queue Q;
    bool visited[mySIZE];
    //for statistical purpose we do not need to store the path
    int source;

public:
    int sIndex;

    __device__
    void start(int s){
        Q.init();
    }
}

```



```

        sIndex = s;
        source = state[s];
        if(&visited == NULL){
            printf("%d: malloc failed for path",sIndex);
            assert(0);
            return;
        }
        for (int i = 0; i < mySIZE; ++i) {
            visited[i] = false;//-1;
        }
        bool fin=false;
        fin = doBFS();

        if(fin){
            count++;
        }
    }
    __device__
    ~eightPuzzle(){
        clean();
    }

    __device__
    void clean(){
        while(!Q.isEmpty()){
            Q.dequeue();
        }
    }

    __device__
    bool doBFS(){
        int currIndex,index;
        int c=0;
        //__shared__
        char tmp[LEN];

        Q.enqueue(source);
        visited[sIndex] = true; //initial; //mark as visited

        Q.enqueue(-1);    //level marker

        int curr,child;
        while(!Q.isEmpty()){
            curr = Q.dequeue();
            if(curr == -1){    //a level has completed
                c++;
                if(Q.isEmpty()){

```

```

        cost[sIndex] = 0;
        break;
    }
    Q.enqueue(-1);    //set marker for next level
    //level completion = next level children are
already in the Q
        continue;
    }
    currIndex = getHash(curr,tmp);
    if(currIndex == gIndex){    //Reached Goal
        cost[sIndex] = c;
        return true;
    }

    //for each child enqueue and mark
    for (int i = UP; i <= RIGHT; ++i) {
        child = move(curr,(Move)i,tmp);
        if(child == -1){
            continue;
        }
        index = getHash(child,tmp);
        if(!visited[index]){
            visited[index] = true;
            Q.enqueue(child);
        }
    }
    }
    return false;
}

};

__device__ eightPuzzle *game;

__global__
void dummy(eightPuzzle *g){
    game = g;
    if(game == NULL)
        printf("Abort");
}

__device__ int iteration;

__global__
void compute() {
    int tid = blockIdx.x*blockDim.x + threadIdx.x;
    game[tid].start(tid+iteration);
}

```

```

__global__ void nextI(){
    iteration += 1008;
}

int main(int argc, char **argv) {
    system("date");
    clock_t t;
    t=clock();

    storeFact<<<1,1>>>();
    cout<<"factorials stored"<<endl;

    populate<<<1,1>>>();
    CUDA_CHECK_RETURN(cudaDeviceSynchronize());
    cout<<sizeof(eightPuzzle)<<endl;
    cout<<sizeof(Queue)<<endl;

    cout<<"alloc done"<<endl;

    cout<<(float(clock()-t))/CLOCKS_PER_SEC<<endl;

    setGoal<<<1,1>>>(123456780);
    CUDA_CHECK_RETURN(cudaDeviceSynchronize());
    cout<<"goal set"<<endl;
    cout<<(float(clock()-t))/CLOCKS_PER_SEC<<endl;

    eightPuzzle *game;
    cudaMalloc((void**)&game,sizeof(eightPuzzle)*1008);

    dummy<<<1,1>>>(game);
    CUDA_CHECK_RETURN(cudaDeviceSynchronize());
    cout<<"game ready"<<endl;
    cout<<(float(clock()-t))/CLOCKS_PER_SEC<<endl;

    int loop = mySIZE/(1008); //9*8*7 = 504
    for (int i = 0; i < loop; ++i) {
        compute<<<1,1008>>>();
        CUDA_CHECK_RETURN(cudaDeviceSynchronize());
        nextI<<<1,1>>>();
        //cout<<(float(clock()-t))/CLOCKS_PER_SEC<<endl;
        //cout<<"loop:"<< i <<endl;
    }
    CUDA_CHECK_RETURN(cudaDeviceSynchronize());
    int c;
    CUDA_CHECK_RETURN(cudaMemcpyFromSymbol(&c,count,sizeof(int),0,cudaMemcpyDe
viceToHost));
    cout<<c<<endl;
    cout<<(float(clock()-t))/CLOCKS_PER_SEC<<endl;

```

```

system("date");

cout<<"===== "<<endl;
int moves[mySIZE];
cudaMemcpyFromSymbol((void*)&moves,cost,sizeof(int)*mySIZE,0,cudaMemcpyDe
viceToHost);
CUDA_CHECK_RETURN(cudaDeviceSynchronize());
cout<<"result copied to host!"<<endl;
cout<<"dumping result"<<endl;
for (int i = 0; i < mySIZE; ++i) {
    cout<<i<<": "<<moves[i]<<endl;
}
}

```